

Raghunandan Kumar

+91 9555808473 · raghu9555k@gmail.com · [LinkedIn](#) · [GitHub](#)

PROFESSIONAL SUMMARY

Full Stack & Agentic AI Engineer with hands-on experience building production-grade LLM pipelines, RAG systems, and multiagent workflows using LangGraph and LangChainStrong foundation in Next.js, Node.js, TypeScript, and Supabase/PostgreSQL. Experienced in integrating AI providers with semantic memory architectures and real-time interfaces.

EXPERIENCE

Ai Engineer and Trainer | Outlier

June 2026 - Present Remote · San Francisco, USA

- Evaluated and rated outputs from large language models (LLMs) based on quality, accuracy, and safety guidelines.
- Trained AI models through structured feedback, improving reasoning, coherence, and response quality.
- Performed data annotation and labelling to support AI model training and fine-tuning datasets.
- Wrote and tested prompts to assess model behaviour and edge-case handling across scenarios.

Software Engineering Intern (AI) | Turbo ML

Apr 2025 – May 2026 · Remote · CA, USA

- Architected an agentic reminder system using Redis Sorted Sets and a custom daemon for time-aware task execution, enabling sub-second scheduling under 24-hour workflow constraints.
- Built multi-tool AI agent integrations connecting Swiggy (food delivery & dine-out), Blinkit (grocery delivery), and Google APIs — enabling dynamic tool selection and user-choice routing.
- Implemented WhatsApp Business API command layer (/help, /reset, /new) with runtime-configurable bot initialization via environment variables, reducing deployment friction.

TECHNICAL SKILLS

Languages: TypeScript, JavaScript, Python, C++, SQL, Golang

Frontend: Next.js, React, Tailwind CSS, CodeMirror 6, HTML5 Canvas

Backend: Node.js, Express.js, FastAPI, Django, WebSocket, REST APIs, WhatsApp Business API

Agentic AI / LLM: LangChain, LangGraph, RAG Pipelines, Prompt Engineering, LLM API Integration Firecrawl, Inngest

Databases: PostgreSQL, Supabase (pgvector), MongoDB, Redis, Convex

Tools: Git, Jira, PyCharm, Docker, Azure (Blob Storage), AWS, GCP, Clerk, Sentry

PROJECTS

Mindly AI — Persistent Memory Agent | Next.js 15, React 19, Groq, pgVector, Supabase, Redis, OAuth 2.0 [Live](#) [Github](#) [Demo-Video](#)

- Engineered a dual-channel memory engine combining pgvector semantic search and PostgreSQL knowledge graphs for cross-session, context-aware AI conversations.
- Built a 4-stage embedding pipeline with automatic API failover (Groq → Gemini → OpenRouter → Ollama) for zerodowntime serverless environments.
- Designed a unified AI provider abstraction layer and background thread execution for relationship extraction — reducing perceived client latency.

AI Interview Coach | Next.js 16, Vapi AI SDK, Groq (Llama 3.3), Firebase, Face-API.js, TypeScript

[Live](#) [Github](#) [Demo-Video](#)

- Integrated Vapi WebSocket SDK with Next.js for low-latency real-time voice interview simulation with live transcription.
- Built client-side face-tracking pipeline using Face-API.js to measure confidence and focus metrics during interviews.
- Created backend evaluation pipeline using Vercel AI SDK, Groq, and Zod to produce structured, schema-validated feedback from LLM output.

Cognito — AI-Powered Browser IDE | Next.js, WebContainer API, Claude Sonnet 4.6, TypeScript, Convex [Live](#) [Github](#)

- Built a browser-based full-stack IDE using WebContainer API and xterm.js — running live Node.js runtimes entirely clientside, zero server provisioning.
- Engineered custom CodeMirror 6 extensions for real-time ghost-text LLM autocomplete and context-aware Cmd+K refactoring via Claude Sonnet 4.6.

EDUCATION